

Computational Geometry

4. Range search

- Data base with names, dates of birth and salary of employees of a company.
- Query: All employees born between 1950 and 1955 earning between 3.000 \$ and 4.000 \$.
- Solution: Compute a single number from the birthday:


$$10.000 \times Year + 100 \times Month + Day$$

- This yields an order for the birthdays that can be used with the income in a *2D-range-query*.

4.1 Motivation

4. Range search

- Geometric interpretation



- Input: Set $P = \{p_1, \dots, p_n\}$ of n points in a k -dimensional space (kd -space) and a k -dimensional *axis-aligned query range* D .
- Output: All points in $P \cap D$.
- Goal
 - Many queries for the same point set P with different query ranges D should be computed as fast as possible.
- Idea
 - Store P in a data structure, that supports range-queries as efficiently as possible.

More examples for different dimensions

- $k = 1$
 - Enumerate all elements in a sequence of keys between a and b .
- $k = 2$
 - Enumerate all cities in a square of 100 km edge length and center in Kaiserslautern.
- $k \geq 3$
 - Data bases queries, e.g. find all persons, that
 - are 20 to 30 years old,
 - earn 30.000 € to 40.000 € and
 - do not posses a cell phone.

Naïve approach

- Test all points sequentially, if the actual point lies in the query range $D = [x_l, x_r] \times [y_l, y_r]$.
- Run time: $O(n)$ for every range search.
- **Inefficient**, if the search result for every range-query contains only a small constant number of points.
- An output-size sensitive algorithm would be beneficial, where the run time depends on the number of points in the query range.

Quad-tree approach (grid-method)

- Place regular grid over the set of points P .
- Consider only points in grid cells, that intersect the query range.
- **Efficient**, if the points are *uniformly* distributed.
- **Inefficient**, if the points are concentrated in one grid area and the grid has many empty grid cells.

BSP-tree approach



- Partition the point set P along arbitrary (hyper-)planes.
- **Inefficient**, because the intersection of a rectangular query range with arbitrary half-spaces is an arbitrary convex polygon (convex polytope in k d).
- **Efficient**, if the (hyper-)planes are axis-aligned:
two-dimensional trees (kd-trees).

- In the sequel only two-dimensional range search, i.e. $k = 2$.
- Extension to higher dimensions is then straight-forward.
- Idea
 - Partition 2d search space in a way similar to a binary search tree for the one-dimensional search space.
 - Use alternating the x - und y -coordinates as keys.
 - Construct a binary tree, representing a partition of the plane:
 - The knots correspond to the n points in 2d.
 - On an even level of the tree (level=depth+1) use the x -coordinate as key, otherwise the y -coordinate.

4.3 *kd*-trees

4. Range search

- A 2d-search tree and the resulting partition of the plane.






Properties

- For a knot in the tree v_x representing point p_x , whose key is its x -coordinate (even level), we have:
 - The left sub-tree of v_x contains the points left of p_x .
 - The right sub-tree of v_x contains the points right of p_x .
 - The sub-trees of children of v_x are determined by y -coordinates.
- For a knot in the tree v_y representing point p_y , whose key is its y -coordinate (odd level), we have:
 - The left sub-tree of v_y contains the points below of p_y .
 - The right sub-tree of v_y contains the points above of p_y .
 - The sub-trees of children of v_y are determined by x -coordinates.

Search in a 2d-tree

- The search is analogous to the usual search in a binary search tree.
- **Attention:** For every visited knot in the tree, the sub-trees of the knot are determined by the x -coordinate **or** the y -coordinate.
- **Attention:** For a range search both sub-trees must be considered.
- Example:



4.3 kd-trees

4. Range search

Algorithm 1: RangeSearch(knot k , direction d , query range D)



```
1: if ( $k \neq \text{NULL}$ ) then {
2:   if ( $d == \text{vertical}$ ) then {
3:      $(l, r) = (D.y_1, D.y_2)$ ;
4:      $\text{coord} = k.y$ ;
5:      $d_{\text{New}} = \text{horizontal}$ ;
6:   } else {
7:      $(l, r) = (D.x_1, D.x_2)$ ;
8:      $\text{coord} = k.x$ ;
9:      $d_{\text{New}} = \text{vertical}$ ;
10:  }
11:  if ( $k \in D$ ) then Add  $k$  to the output;
12:  if ( $l < \text{coord}$ ) then RangeSearch( $k.\text{left}$ ,  $d_{\text{New}}$ ,  $D$ );
13:  if ( $r > \text{coord}$ ) then RangeSearch( $k.\text{right}$ ,  $d_{\text{New}}$ ,  $D$ );
14: }
```

Call: RangeSearch($1, n, \text{root}, \text{vertical}$)

4.3 *kd*-trees

4. Range search

- Example



Run time

- If the 2d-tree is balanced, the height of the tree is $O(\log n)$.
- If the 2d-tree is unbalanced, a path in the tree might degenerate to a linear list resulting in a bad run time.
- Goal: The construction of the tree must guarantee the balanced-ness.
- Solution: Partition the points at the **median** of the sequence of x - and y -coordinates.




Partitioning at the median

- Sort all points both by x - and y -coordinates
 - ➔ Two sorted sequences X and Y .
- Subdivide Y at its median and take it as the root of the tree
 - ➔ Two sub-sequences Y_1 and Y_2 .
- Partition X into two sequences X_1 and X_2 , such that X_1 contains the same points as Y_1 and X_2 the same points as Y_2 .
- Partition X_1 and X_2 recursively at the median and subdivide Y_1 and Y_2 accordingly as above, until these sequences contain only one point.
 - These points are the leafs of the tree.

4.3 *kd*-trees

4. Range search

■ Example for the partition



4.3 kd-trees

4. Range search

- Global, pre-sorted sequences X and Y .

Algorithm 2: ConstructBalanced2DTree

Input: left index l , right index r , knot k , direction d

```
1: if ( $l \leq r$ ) then {  
2:    $m = \lfloor \frac{l+r}{2} \rfloor$ ;  
3:   if ( $d == \text{vertical}$ ) then {  
4:      $k.\text{value} = Y[m]$ ;  
5:     PartitionField( $X, l, r, m$ );  
6:   } else {  
7:      $k.\text{value} = X[m]$ ;  
8:     PartitionField( $Y, l, r, m$ );  
9:   }  
10:  ConstructBalanced2dTree( $l, m-1, k.\text{left}, !d$ );  
11:  ConstructBalanced2dTree( $m+1, r, k.\text{right}, !d$ );  
12: }
```

Call: ConstructBalanced2DTree($1, n, \text{root}, \text{vertical}$)

Run time

■ Construction

- Sorting of sequences: $O(n \log n)$.
- Partitioning of the sequences: $O(n)$.
- Recursive calls: $T(n) = T\left(\left\lceil \frac{n-1}{2} \right\rceil\right) + T\left(\left\lfloor \frac{n-1}{2} \right\rfloor\right) + O(n)$
- Solution of the recursion: $O(n \log n)$.

■ Range search in a balanced 2d-tree

- Run time for R points in the query range D : $O(\sqrt{n} + R)$.
- Proof: Black board.
- For small R this is much faster than testing all points.

Higher dimensions

- Use the same algorithm!
- Partition for the construction of the tree cycle through the dimensions one by one.
- For the search also cycle through the dimensions one after the other.
- Run time and storage in k dimensions
 - Construction: $O(kn \log n)$
 - Range search: $O(kn^{1-\frac{1}{k}} + R)$
 - Storage: $O(n)$ (A *kd*-tree is a binary tree with n leaves.)

- For large n and small R the query time of a kd -tree is relatively large.
- Principle of kd -trees:
Comparisons of x - and y -coordinates alternate.
- Principle of range-trees:
Do comparisons of x - and y -coordinates subsequently.

- First the x -coordinates:
 - First level data structure: Use a balanced binary tree for the x -coordinates of the points.
 - Query: Use a **one-dimensional range query** to find all points with x -coordinates in $[x_l, x_r]$.
 - This yields a list of candidates that lie potentially in the query range D .
- Second the y -coordinates:
 - Second level data structure: Use a balanced binary tree for the y -coordinates of the candidates.
 - Query: Use a **one-dimensional range query** to find all candidates with y -coordinates in $[y_l, y_r]$.

4.4.1 1d-range-search

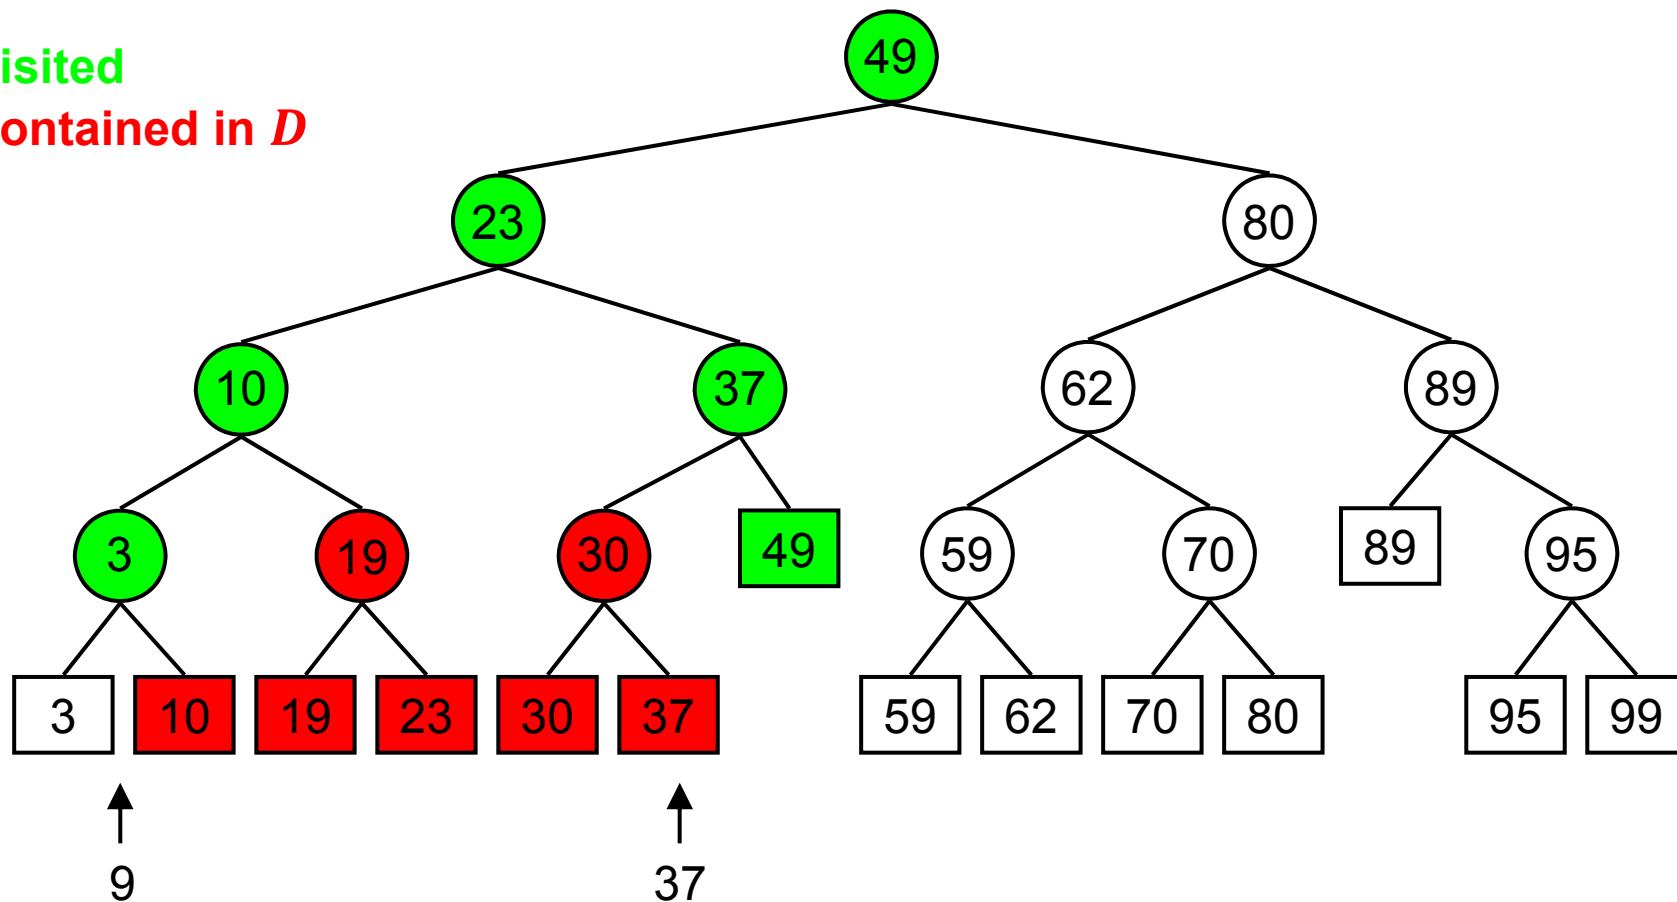
4. Range search

4.4 Range-trees

Example (1d-range-query): Search for $D = [9, 37]$

visited

contained in D



1D-Range-Search: Search for all leaves in a balanced, binary search tree whose key lies in $[x_l, x_r]$.

Idea: Search for x_l and x_r in the tree simultaneously:

1. Find the node v_{split} where the paths to x_l and x_r split.
2. From v_{split} follow the path to x_l .
 - At a node where this path turns left, report all leaves of the right sub-tree.
3. From v_{split} follow the path to x_r .
 - At a node where this path turns right, report all leaves of the left sub-tree.

4.4.1 1d-range-search

4. Range search

4.4 Range-trees

Algorithm 3: FindSplitNode(tree T , left bound x_l , right bound x_r)

Output: Node where paths to x_l and x_r splits, or the leaf where both paths end.

```
1:  $v = \text{Root}(T);$ 
2: while ( $v$  is not a leaf and  $(x_l > x_v \text{ or } x_r \leq x_v)$ ) do {
3:   if  $(x_r \leq x_v)$  then  $v = \text{LeftChild}(v);$ 
4:   else  $v = \text{RightChild}(v);$ 
5: }
6: return  $v;$ 
```

4.4.1 1d-range-search

4. Range search

4.4 Range-trees

Algorithm 4: 1dRangeSearch(tree T , left bound x_l , right bound x_r)

Output: All point in T that lie in the interval $[x_l, x_r]$

```
1:  $v_{\text{split}} = \text{FindSplitNode}(T, x_l, x_r);$ 
2: if ( $v_{\text{split}}$  is a leaf) then Return the point in  $v_{\text{split}}$ ,
   if necessary;
3: else {
4:    $v = \text{LeftChild}(v_{\text{split}});$ 
5:   while ( $v$  is not a leaf) do {
6:     if ( $x_l \leq x_v$ ) then {
7:       ReportSubtree(RightChild( $v$ ));
8:        $v = \text{LeftChild}(v);$ 
9:     } else  $v = \text{RightChild}(v);$ 
10:  }
11:  Return the point in  $v$ , if necessary;
12:  // Analogous for RightChild( $v_{\text{split}}$ ) and  $x_r$ ;
  :
  : }
```

Proposition 1

A one-dimensional range search can be computed using $O(n)$ storage in $O(n \log n)$ for the pre-processing and $O(R + \log n)$ for the range search, where R is the number of points, that lie in the query range.

- Proof:**
- **Pre-processing:** Store the 1d-point sequence in a balanced, binary search tree (e.g. AVL-tree)
 - Storage: $O(n)$
 - Construction time: $O(n \log n)$
 - **Range query:**
 - To compute v_{split} and to follow the path to x_l and x_r takes $O(\log n)$, because the time spend in each node is $O(1)$ and the tree is balanced.
 - ReportSubtree is linear in the reported points which is in total $O(R)$.
 - Total: $O(R + \log n)$.


2D-Range-Search: Search for all leaves in a range tree whose key lies in $[x_l, x_r] \times [y_l, y_r]$.

- First level data structure: Balanced, binary search tree for the x -coordinates of the points.
- Every sub-tree rooted at a node v represents a subset of the points, the so-called **canonical subset** $P(v)$.
- For the range query we need only those points in the canonical subsets whose y -coordinates are in $[y_l, y_r]$.
- Second level data structure: Balanced, binary search tree for the y -coordinates of the points in the canonical subsets.

Range-tree

- Balanced, binary search tree for the x -coordinates of the points.
- At every inner node v store a balanced, binary search tree for the y -coordinates of the points in $P(v)$.

- This is the so-called **associated structure** $T_{\text{assoc}}(v)$.
- The leaves of $T_{\text{assoc}}(v)$ hold the points.



Proposition 2

A range tree for n 2d points uses $O(n \log n)$ meomory and can be constructed in $O(n \log n)$.

Proof:

■ Storage:

- At each level every point is stored in exactly one associated structure, which uses linear memory.
- Thus, at every level $O(n)$ meomory is used in total for all associated structures.
- Because there are $O(\log n)$ levels, $O(n \log n)$ of memory is used.

■ Construction

- To build a range-tree pre-compute two lists of the points sorted by x - and y -coordinates.
- From these two lists build the search trees bottom up in $O(n \log n)$.

4.4 Range-trees

4. Range search

Algorithm 5: 2dRangeSearch(tree T , query range $[x_l, x_r] \times [y_l, y_r]$)

Output: All point in T that lie in the query range $[x_l, x_r] \times [y_l, y_r]$

```
1:  $v_{\text{split}} = \text{FindSplitNode}(T, x_l, x_r);$ 
2: if ( $v_{\text{split}}$  is a leaf) then Return the point in  $v_{\text{split}}$ ,  
   if necessary;
3: else {
4:    $v = \text{LeftChild}(v_{\text{split}});$ 
5:   while ( $v$  is not a leaf) do {
6:     if ( $x_l \leq x_v$ ) then {
7:       1dRangeSearch( $T_{\text{assoc}}(\text{RightChild}(v)), y_l, y_r$ );
8:        $v = \text{LeftChild}(v);$ 
9:     } else  $v = \text{RightChild}(v);$ 
10:  }
11:  Return the point in  $v$ , if necessary;
12:  // Analogous for RightChild( $v_{\text{split}}$ ) and  $x_r$ ;
  :
  : }
```

Proposition 3

A query in a range tree of n 2d points takes $O(R + \log^2 n)$ time, where R is the number of points, that lie in the query range.

Proof:

- Every call of 1DRangeSearch at a node v takes $O(R_v + \log n)$, where R_v is the number of points in $D \cap P(v)$, i.e. $\sum R_v = R$.
- Hence, the total time spend is in 1DRangeSearch
$$\sum O(R_v + \log n),$$

where the summation is over all nodes v that are visited on the search path to x_l and x_r .

- Because the search paths to x_l and x_r have length $O(\log n)$, the total run time is $O(R + \log^2 n)$.

Higher dimensions

- Construct a k -dimensional range-tree recursively, i.e. the second level data structure is a $(k - 1)$ -dimensional range-tree.
- Run time and storage in k dimensions
 - Construction: $O(n \log^{k-1} n)$
 - Range search: $O(R + \log^k n)$
 - Storage: $O(n \log^{k-1} n)$